

Mixture Models: Model-Based Clustering

Emily Fox & Carlos Guestrin


Machine Learning Specialization

University of Washington

Why a probabilistic approach?

Learn user preferences

Set of clustered documents read by user



Cluster 1



Cluster 2



Cluster 3


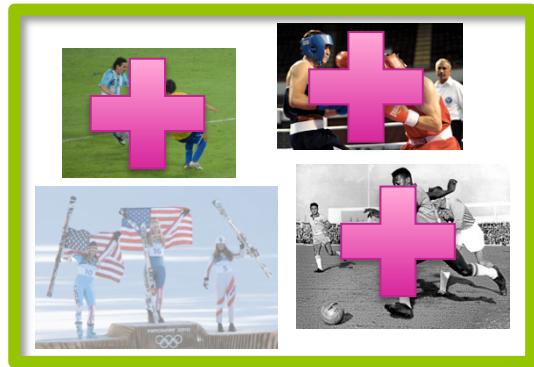


Cluster 4



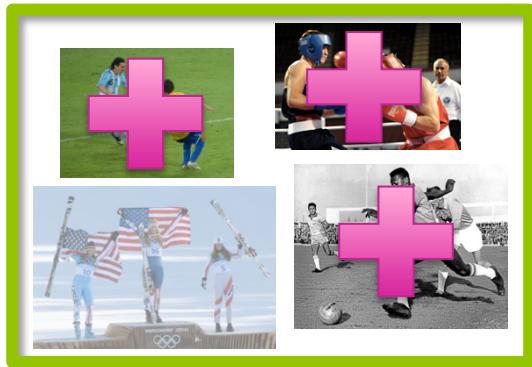
Use feedback
to learn user
preferences
over topics

Uncertainty in cluster assignments



Slightly closer to Cluster 4 than Cluster 2, but count fully for Cluster 4?

Uncertainty in cluster assignments



Hard assignments
don't tell full story

Other limitations of k-means

Assign observations to closest cluster center



$$z_i \leftarrow \arg \min_j \| \mu_j - \mathbf{x}_i \|_2^2$$

Only center matters




Can use weighted Euclidean,
but requires *known* weights

Still assumes all clusters have
the same axis-aligned ellipses

Equivalent to assuming
spherically symmetric clusters



Failure modes of k-means



Motivates probabilistic model: mixture model


- Provides **soft assignments** of observations to clusters (uncertainty in assignment)
 - e.g., 54% chance document is **world news**, 45% **science**, 1% **sports**, and 0% **entertainment**
- Accounts for cluster **shapes** not just **centers**
- Enables **learning weightings** of dimensions
 - e.g., how much to weight each word in the vocabulary when computing cluster assignment

Mixture models

Motivating application: Clustering images

Discover groups of similar images

- Ocean
- Pink flower
- Dog
- Sunset
- Clouds
- ...




Simple image representation

Consider average red, green, blue pixel intensities



[R = 0.05, G = 0.7, B = 0.9]




[R = 0.85, G = 0.05, B = 0.35]



[R = 0.02, G = 0.95, B = 0.4]

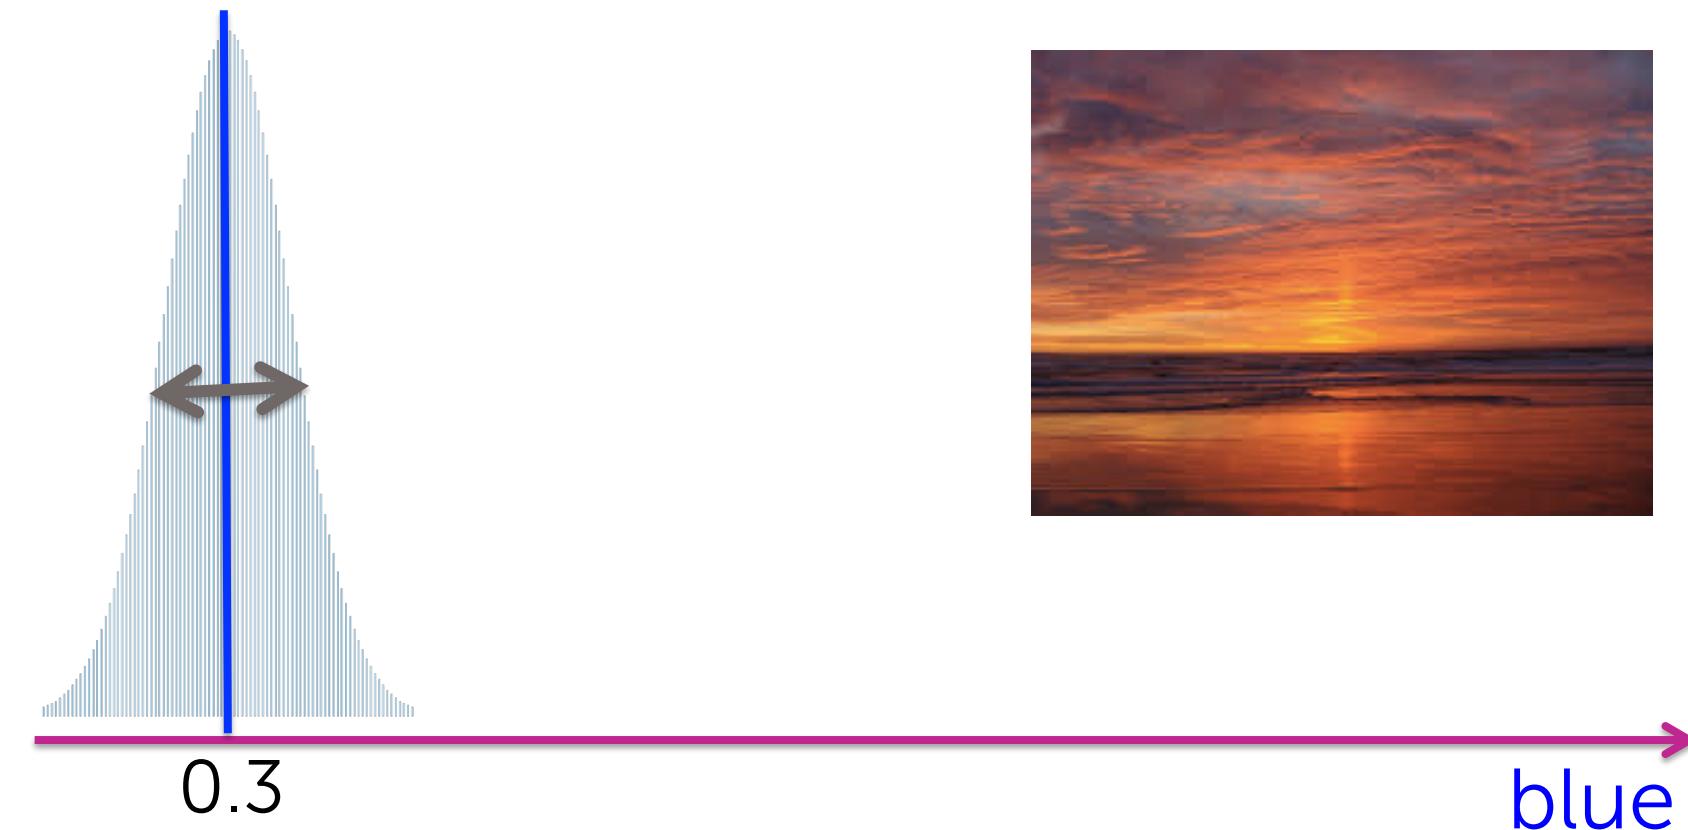
Distribution over all **cloud** images

Let's look at just the **blue** dimension




Distribution over all **sunset** images

Let's look at just the **blue** dimension




Distribution over all forest images

Let's look at just the **blue** dimension

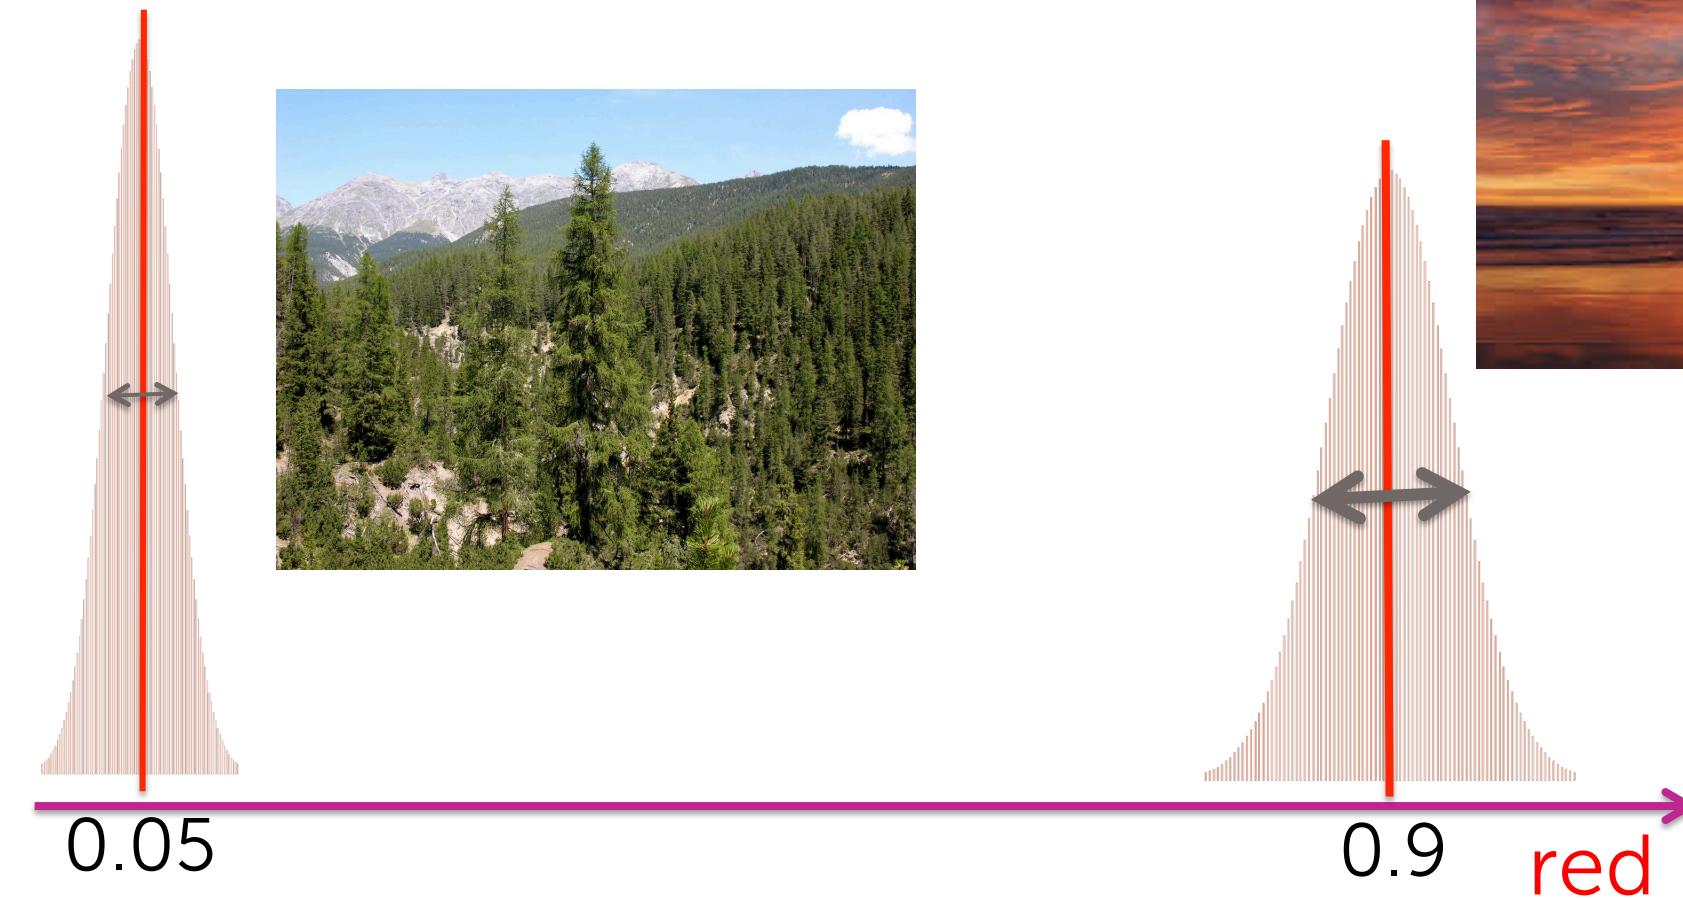


Distribution over all images



Can be distinguished along other dim


Now look at the **red** dimension



Background: Gaussian distributions


Model for a given image type

For **each dimension** of the [R, G, B] vector,
and **each image type**, assume a
Gaussian distribution over color intensity




1D Gaussians

Fully specified by **mean** μ and **variance** σ^2
(or **standard deviation** σ)




Notating a 1D Gaussian distribution

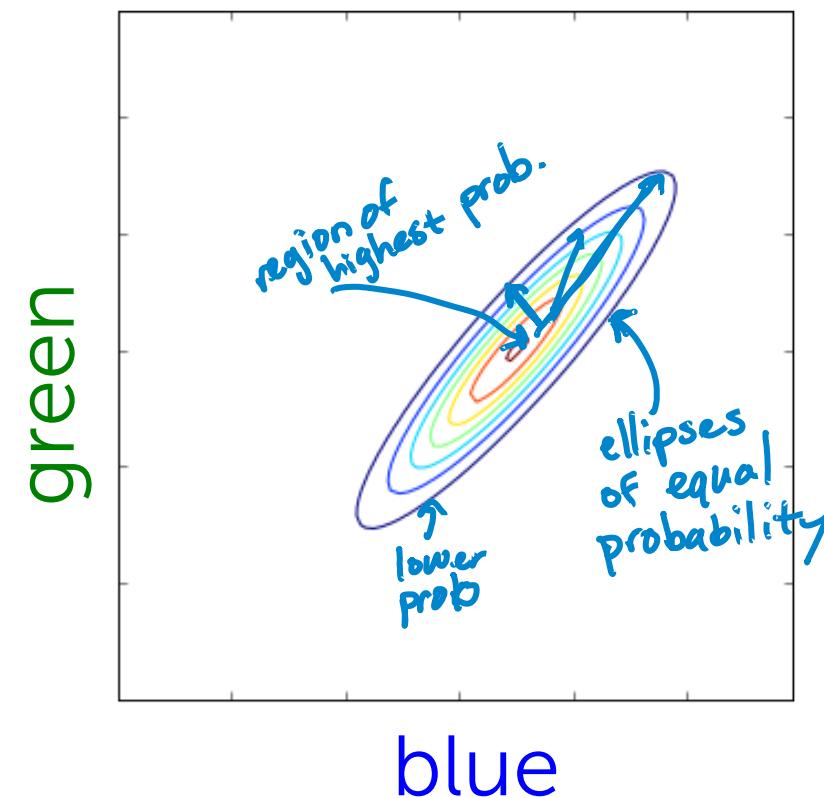


2D Gaussians – Bird's eye view

3D mesh plot



Contour plot




2D Gaussians – Parameters

Fully specified by **mean μ** and **covariance Σ**

$$\mu = [\mu_{\text{blue}}, \mu_{\text{green}}]$$

mean centers the distribution in 2D



2D Gaussians – Parameters

Fully specified by **mean** μ and **covariance** Σ

$$\mu = [\mu_{\text{blue}}, \mu_{\text{green}}]$$


$$\Sigma = \begin{pmatrix} \sigma_{\text{blue}}^2 & \sigma_{\text{blue},\text{green}} \\ \sigma_{\text{green},\text{blue}} & \sigma_{\text{green}}^2 \end{pmatrix}$$

covariance determines
orientation + spread

green


Σ

blue




Covariance structures





$$\Sigma = \begin{pmatrix} \sigma^2 & 0 \\ 0 & \sigma^2 \end{pmatrix}$$



$$\Sigma = \begin{pmatrix} \sigma_B^2 & 0 \\ 0 & \sigma_G^2 \end{pmatrix}$$




$$\Sigma = \begin{pmatrix} \sigma_B^2 & \sigma_{B,G} \\ \sigma_{G,B} & \sigma_G^2 \end{pmatrix}$$



Notating a multivariate Gaussian



$$N(\mathbf{x} \mid \boldsymbol{\mu}, \boldsymbol{\Sigma})$$

Random **vector**
e.g., [R, G, B] intensities




Mixture of Gaussians

Model as Gaussian per category/cluster



Jumble of unlabeled images



HISTOGRAM



How do we model
this distribution?





Model of jumble of unlabeled images



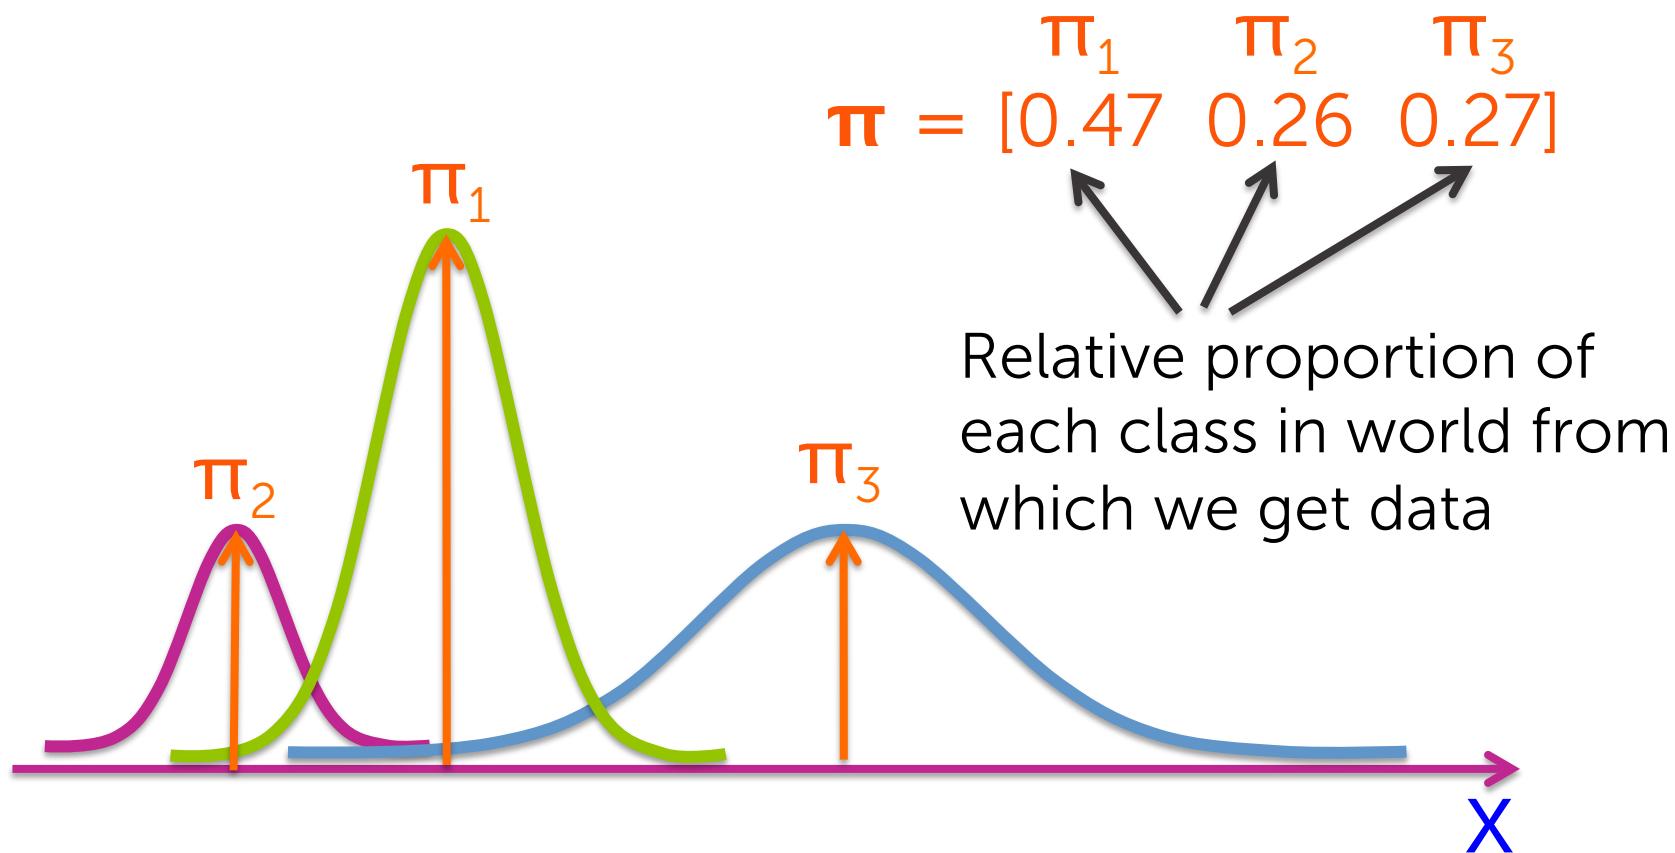
What if image types not equally represented?

e.g., forest images are very likely in the collection



Combination of weighted Gaussians


Associate a weight π_k with each Gaussian component



Combination of weighted Gaussians


Associate a weight π_k with each Gaussian component

$$\boldsymbol{\pi} = [\pi_1 \quad \pi_2 \quad \pi_3] = [0.47 \quad 0.26 \quad 0.27]$$




Mixture of Gaussians (1D)

Each mixture component represents a unique cluster specified by:



Mixture of Gaussians (general)



Each mixture component represents a unique cluster specified by:

$$\{\pi_k, \mu_k, \Sigma_k\}$$

According to the model...


Without observing the image content, what's the probability it's from cluster k? (e.g., prob. of seeing "clouds" image)

$$p(z_i = k) = \underline{\pi_k} \quad \text{prior}$$

cluster assignment for obs. x_i

Given observation x_i is from cluster k, what's the likelihood of seeing x_i ? (e.g., just look at distribution for "clouds")

$$p(x_i | z_i = k, \mu_k, \Sigma_k) = N(x_i | \mu_k, \Sigma_k) \quad \text{likelihood}$$



Document clustering

Discover groups of related documents




Document representation



$x_i =$ tf-idf vector


Mixture of Gaussians for clustering documents

Space of all documents
(really lives in \mathbb{R}^V for vocab size V)



Mixture of Gaussians for clustering documents


Space of all documents
(really lives in \mathbb{R}^V for vocab size V)



Make soft assignments
of docs to each
Gaussian

Counting parameters

Each cluster has $\{\pi_k, \mu_k, \Sigma_k\}$



In 2D:

$$\Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{1,2} \\ \sigma_{2,1} & \sigma_2^2 \end{pmatrix}$$


Counting parameters

Each cluster has $\{\pi_k, \mu_k, \Sigma_k\}$

In V (vocab size) dims:


$$\Sigma =$$

$$\frac{V(V+1)}{2}$$



Restricting to diagonal covariance


Each cluster has $\{\pi_k, \mu_k, \Sigma_k \text{ diagonal}\}$



V params

$$\Sigma = \begin{pmatrix} \sigma_1^2 & & & \\ & \sigma_2^2 & & \\ & & \sigma_3^2 & \\ & & & \ddots & \ddots & \ddots \\ & & & & \sigma_V^2 & \end{pmatrix}$$


Restrictive assumption, but...



- Can **learn** weights on dimensions (e.g., weights on words in vocab)
- Can learn **cluster-specific** weights on dimensions

Still more flexible than k-means

Spherically
symmetric clusters



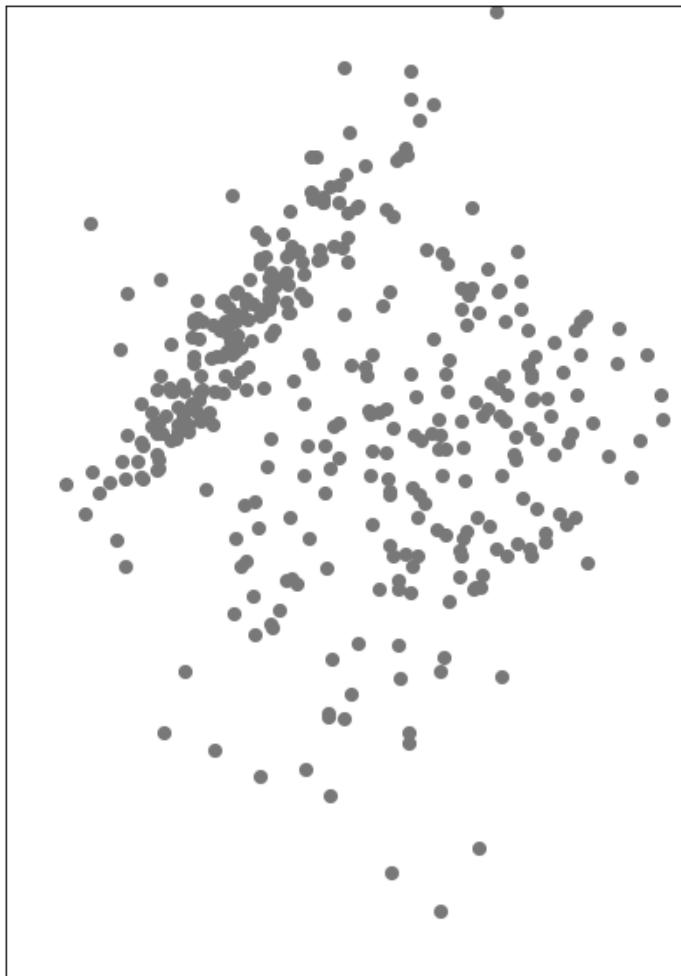
Specify weights...

All clusters have same
axis-aligned ellipses


Inferring soft assignments with expectation maximization (EM)

Inferring cluster labels

Data




Desired soft assignments



Part 1:
What if we knew the cluster
parameters $\{\pi_k, \mu_k, \Sigma_k\}$?

Compute responsibilities



$$r_{ik} = p(z_i = k \mid \{\pi_j, \mu_j, \Sigma_j\}_{j=1}^K, x_i)$$

Responsibility cluster k takes for observation i

r_{ik} is a random variable

probability of assignment to cluster k


$\{\pi_j, \mu_j, \Sigma_j\}_{j=1}^K$ are fixed values defining the distribution

clusters

"given"

given model parameters and observed value


Responsibilities in pictures



Green cluster
takes more
responsibility



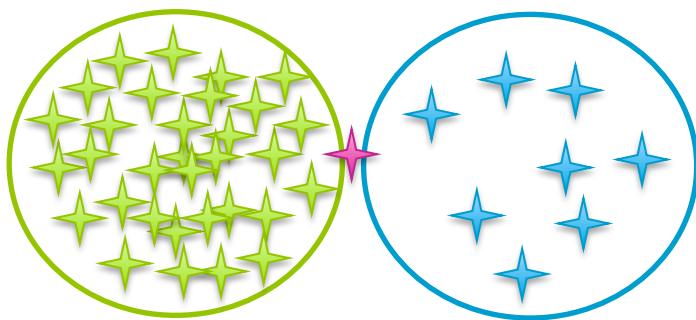
Blue cluster
takes more
responsibility



Uncertain...
split
responsibility


Responsibilities in pictures

Need to weight by cluster probabilities,
not just cluster shapes



Still **uncertain**,
but **green** cluster seems
more probable...
takes more responsibility

Responsibilities in equations




$$r_{ik} = \pi_k N(x_i | \mu_k, \Sigma_k)$$

Initial probability of being from cluster k

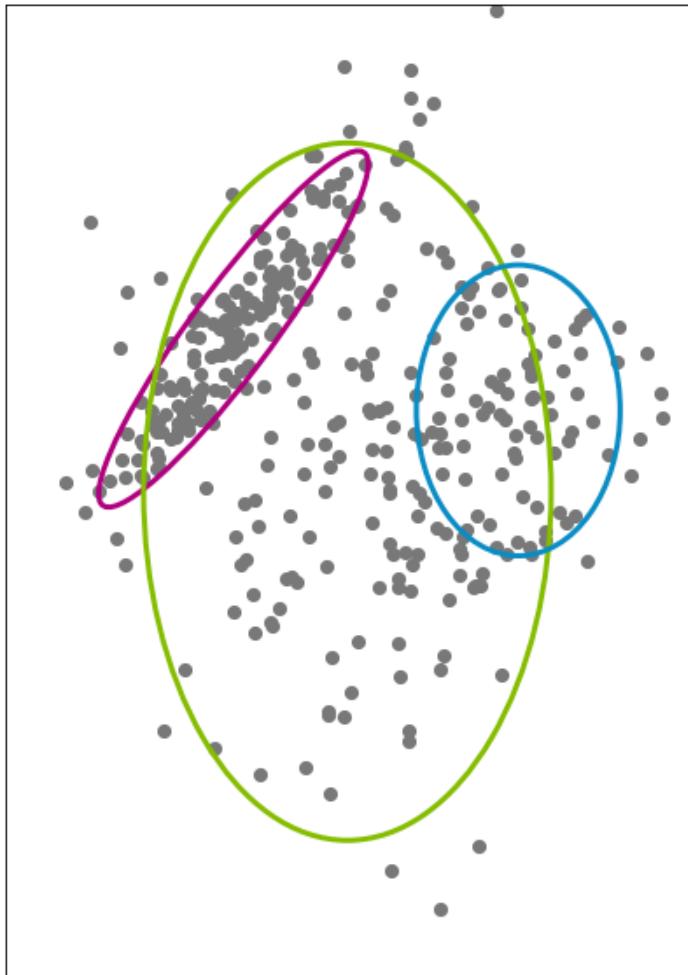
Responsibility cluster k takes for observation i

How likely is the observed value x_i under this cluster assignment?



very unlikely under the green cluster,
even though the prior on green is higher

Responsibilities in equations



$$r_{ik} = \frac{\pi_k N(x_i | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j N(x_i | \mu_j, \Sigma_j)}$$

Responsibility cluster k takes for observation i

Normalized over all possible cluster assignments


Recall: According to the model...

Without observing the image content, what's the probability it's from cluster k? (e.g., prob. of seeing "clouds" image)

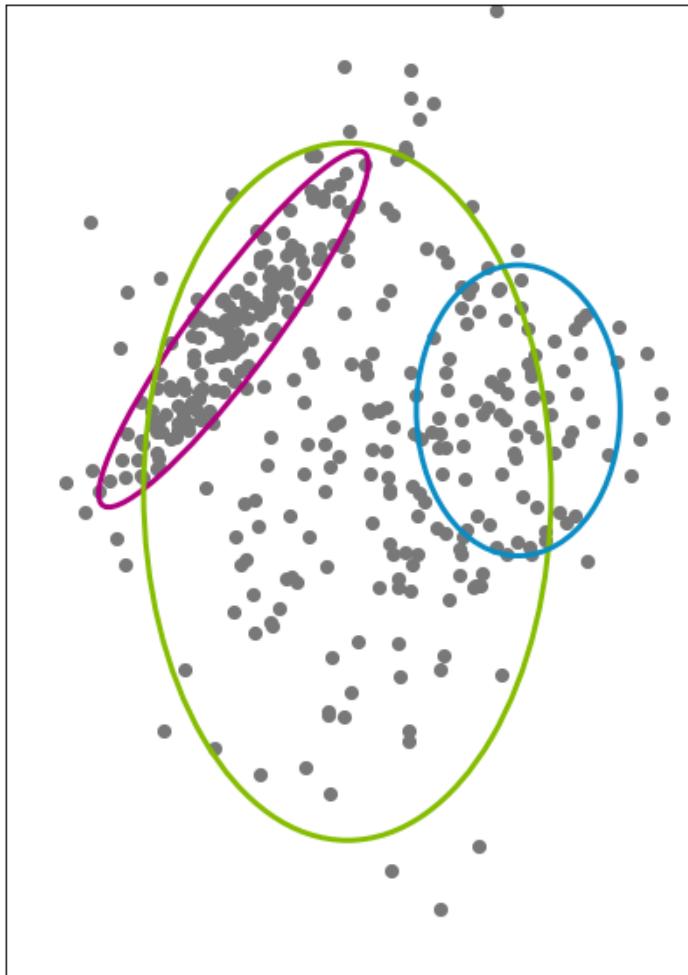
$$p(z_i = k) = \pi_k$$

Given observation \mathbf{x}_i is from cluster k, what's the likelihood of seeing \mathbf{x}_i ? (e.g., just look at distribution for "clouds")

$$p(x_i | z_i = k, \mu_k, \Sigma_k) = N(x_i | \mu_k, \Sigma_k)$$



Part 1 summary



Desired soft assignments (**responsibilities**) are **easy** to compute when cluster parameters $\{\pi_k, \mu_k, \Sigma_k\}$ are known

But, we don't know these!

Responsibility calculation as and application of Bayes' rule

OPTIONAL



An application of Bayes' rule

Responsibility cluster k takes for observation i

$$r_{ik} = p(z_i \in k \mid \{\pi_j, \mu_j, \Sigma_j\}_{j=1}^K, \mathcal{B}_i)$$
$$= \frac{\pi_k N(x_i \mid \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j N(x_i \mid \mu_j, \Sigma_j)}$$

$p(z_i \in k \mid \text{params})$

$p(\mathcal{B}_i \mid z_i \in k, \text{params})$

An application of Bayes' rule

Responsibility cluster k takes for observation i

$$r_{ik} = p(z_i \in k \mid \{\pi_j, \mu_j, \Sigma_j\}_{j=1}^K, \mathcal{B}_i)$$
$$= \frac{\pi_k N(x_i \mid \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j N(x_i \mid \mu_j, \Sigma_j)}$$

π_k $N(x_i \mid \mu_k, \Sigma_k)$

π_j $N(x_i \mid \mu_j, \Sigma_j)$

$p(z_i \in j \mid \text{params})$ $p(\mathcal{B}_i \mid z_i \in j, \text{params})$

An application of Bayes' rule

$$\underline{r_{ik}} = p(A|B, \text{params})$$

$$= \frac{p(A|\text{params})p(B|A, \text{params})}{\sum_C p(C|\text{params})p(B|C, \text{params})}$$

$\underbrace{p(B, C | \text{params})}_{C}$

$$= \frac{p(A|\text{params})p(B|A, \text{params})}{p(B|\text{params})}$$

more general
form of $p(A|B) = \frac{p(A)p(B|A)}{p(B)}$


events B and C

$\downarrow \downarrow$

$$\sum_C p(B, C | \text{params}) = p(B | \text{params})$$

Part 2a:
Imagine we knew the cluster
(hard) assignments z_i

Estimating cluster parameters



Imagine we know the cluster assignments

Estimation problem decouples across clusters

Is green point informative of fuchsia cluster parameters?

NO!

Data table decoupling over clusters

R	G	B	Cluster
$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$	3
$\mathbf{x}_2[1]$	$\mathbf{x}_2[2]$	$\mathbf{x}_2[3]$	3
$\mathbf{x}_3[1]$	$\mathbf{x}_3[2]$	$\mathbf{x}_3[3]$	3
$\mathbf{x}_4[1]$	$\mathbf{x}_4[2]$	$\mathbf{x}_4[3]$	1
$\mathbf{x}_5[1]$	$\mathbf{x}_5[2]$	$\mathbf{x}_5[3]$	2
$\mathbf{x}_6[1]$	$\mathbf{x}_6[2]$	$\mathbf{x}_6[3]$	2

Maximum likelihood estimation

R	G	B	Cluster
$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$	3
$\mathbf{x}_2[1]$	$\mathbf{x}_2[2]$	$\mathbf{x}_2[3]$	3
$\mathbf{x}_3[1]$	$\mathbf{x}_3[2]$	$\mathbf{x}_3[3]$	3

Estimate $\{\pi_k, \mu_k, \Sigma_k\}$
given data assigned
to cluster k

maximum likelihood estimation
(MLE)

Find parameters that maximize the
score, or *likelihood*, of data

Mean/covariance MLE

Sum characteristics

R	G	B	Cluster
$x_1[1]$	$x_1[2]$	$x_1[3]$	3
$x_2[1]$	$x_2[2]$	$x_2[3]$	3
$x_3[1]$	$x_3[2]$	$x_3[3]$	3

divide by 3
(the total # of obs.)

"best estimate" $\hat{\mu}_k = \frac{1}{N_k} \sum_{i \text{ in } k} x_i$ ← average data points in cluster k

$$\hat{\Sigma}_k = \frac{1}{N_k} \sum_{i \text{ in } k} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T$$

Scalar case: $\hat{\sigma}_k^2 = \frac{1}{N_k} \sum_{i \text{ in } k} (x_i - \hat{\mu}_k)^2$

Cluster proportion MLE

R	G	B	Cluster
$\mathbf{x}_4[1]$	$\mathbf{x}_4[2]$	$\mathbf{x}_4[3]$	1

R	G	B	Cluster
$\mathbf{x}_5[1]$	$\mathbf{x}_5[2]$	$\mathbf{x}_5[3]$	2
$\mathbf{x}_6[1]$	$\mathbf{x}_6[2]$	$\mathbf{x}_6[3]$	2

R	G	B	Cluster
$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$	3
$\mathbf{x}_2[1]$	$\mathbf{x}_2[2]$	$\mathbf{x}_2[3]$	3
$\mathbf{x}_3[1]$	$\mathbf{x}_3[2]$	$\mathbf{x}_3[3]$	3


obs in cluster k

$$\hat{\pi}_k = \frac{N_k}{N}$$

total # of obs

True for general mixtures of i.i.d. data,
not just Gaussian clusters

Part 2a summary




needed to compute soft assignments
Cluster parameters are simple
to compute if we know the
cluster assignments

But, we don't know these!

Part 2b:

What can we do with just soft assignments r_{ij} ?

Estimating cluster parameters from soft assignments



Instead of having a full observation \mathbf{x}_i in cluster k , just allocate a portion r_{ik}

\mathbf{x}_i divided across all clusters,
as determined by r_{ik}

Maximum likelihood estimation from soft assignments

Just like in boosting with weighted observations...

R	G	B	r_{i1}	r_{i2}	r_{i3}
$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$	0.30	0.18	0.52
$\mathbf{x}_2[1]$	$\mathbf{x}_2[2]$	$\mathbf{x}_2[3]$	0.01	0.26	0.73
$\mathbf{x}_3[1]$	$\mathbf{x}_3[2]$	$\mathbf{x}_3[3]$	0.002	0.008	0.99
$\mathbf{x}_4[1]$	$\mathbf{x}_4[2]$	$\mathbf{x}_4[3]$	0.75	0.10	0.15
$\mathbf{x}_5[1]$	$\mathbf{x}_5[2]$	$\mathbf{x}_5[3]$	0.05	0.93	0.02
$\mathbf{x}_6[1]$	$\mathbf{x}_6[2]$	$\mathbf{x}_6[3]$	0.13	0.86	0.01



52% chance
this obs is in
cluster 3

Total weight in cluster:
(effective # of obs)

1.242 2.8 2.42

Maximum likelihood estimation from soft assignments

R	G	B	r_{i1}	r_{i2}	r_{i3}
$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$	0.30	0.18	0.52
$\mathbf{x}_2[1]$	$\mathbf{x}_2[2]$	$\mathbf{x}_2[3]$	0.01	0.26	0.73
$\mathbf{x}_3[1]$	$\mathbf{x}_3[2]$	$\mathbf{x}_3[3]$	0.002	0.008	0.99
$\mathbf{x}_4[1]$	$\mathbf{x}_4[2]$	$\mathbf{x}_4[3]$	0.75	0.10	0.15
$\mathbf{x}_5[1]$	$\mathbf{x}_5[2]$	$\mathbf{x}_5[3]$	0.05	0.93	0.02
$\mathbf{x}_6[1]$	$\mathbf{x}_6[2]$	$\mathbf{x}_6[3]$	0.13	0.86	0.01

Maximum likelihood estimation from soft assignments

R	G	B	Cluster 1 weights
$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$	0.30
$\mathbf{x}_2[1]$	R	G	Cluster 2 weights
$\mathbf{x}_3[1]$			
$\mathbf{x}_4[1]$	$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$
$\mathbf{x}_5[1]$	$\mathbf{x}_2[1]$	R	Cluster 3 weights
$\mathbf{x}_6[1]$	$\mathbf{x}_3[1]$	G	
$\mathbf{x}_4[1]$	$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$
$\mathbf{x}_5[1]$	$\mathbf{x}_2[1]$	$\mathbf{x}_2[2]$	$\mathbf{x}_2[3]$
$\mathbf{x}_6[1]$	$\mathbf{x}_3[1]$	$\mathbf{x}_3[2]$	$\mathbf{x}_3[3]$
	$\mathbf{x}_4[1]$	$\mathbf{x}_4[2]$	$\mathbf{x}_4[3]$
	$\mathbf{x}_5[1]$	$\mathbf{x}_5[2]$	$\mathbf{x}_5[3]$
	$\mathbf{x}_6[1]$	$\mathbf{x}_6[2]$	$\mathbf{x}_6[3]$

Cluster-specific location/shape MLE

R	G	B	Cluster 1 weights
$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$	0.30
$\mathbf{x}_2[1]$	$\mathbf{x}_2[2]$	$\mathbf{x}_2[3]$	0.01
$\mathbf{x}_3[1]$	$\mathbf{x}_3[2]$	$\mathbf{x}_3[3]$	0.002
$\mathbf{x}_4[1]$	$\mathbf{x}_4[2]$	$\mathbf{x}_4[3]$	0.75
$\mathbf{x}_5[1]$	$\mathbf{x}_5[2]$	$\mathbf{x}_5[3]$	0.05
$\mathbf{x}_6[1]$	$\mathbf{x}_6[2]$	$\mathbf{x}_6[3]$	0.13

Compute cluster parameter estimates
with weights on each row operation

$$\hat{\mu}_k = \frac{1}{N_k^{\text{soft}}} \sum_{i=1}^N r_{ik} x_i$$
$$\hat{\Sigma}_k = \frac{1}{N_k^{\text{soft}}} \sum_{i=1}^N r_{ik} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T$$

$$N_k^{\text{soft}} = \sum_{i=1}^N r_{ik}$$

Total weight in cluster k
= effective # obs

MLE of cluster proportions $\hat{\pi}_k$

r_{i1}	r_{i2}	r_{i3}
0.30	0.18	0.52
0.01	0.26	0.73
0.002	0.008	0.99
0.75	0.10	0.15
0.05	0.93	0.02
0.13	0.86	0.01

Total weight
in cluster:

1.242 | 2.8 | 2.42

Total weight
in dataset:

6

datapoints N

$$\hat{\pi}_k = \frac{N_k^{\text{soft}}}{N}$$

$$N_k^{\text{soft}} = \sum_{i=1}^N r_{ik}$$

Total weight in cluster k
= effective # obs

Estimate cluster
proportions from
relative weights

Defaults to hard assignment case when r_{ij} in {0,1}

Hard assignments have:

$$r_{ik} = \begin{cases} 1 & i \text{ in } k \\ 0 & \text{otherwise} \end{cases}$$

R	G	B	r_{i1}	r_{i2}	r_{i3}
$\mathbf{x}_1[1]$	$\mathbf{x}_1[2]$	$\mathbf{x}_1[3]$	0	0	1
$\mathbf{x}_2[1]$	$\mathbf{x}_2[2]$	$\mathbf{x}_2[3]$	0	0	1
$\mathbf{x}_3[1]$	$\mathbf{x}_3[2]$	$\mathbf{x}_3[3]$	0	0	1
$\mathbf{x}_4[1]$	$\mathbf{x}_4[2]$	$\mathbf{x}_4[3]$	1	0	0
$\mathbf{x}_5[1]$	$\mathbf{x}_5[2]$	$\mathbf{x}_5[3]$	0	1	0
$\mathbf{x}_6[1]$	$\mathbf{x}_6[2]$	$\mathbf{x}_6[3]$	0	1	0

One-hot encoding of
cluster assignment

Total weight in cluster:



Equating the estimates...

$$\hat{\pi}_k = \frac{N_k^{\text{Soft}}}{N}$$

$$N_k^{\text{soft}} = \sum_{i=1}^N r_{ik}$$

if $r_{ik} > 0, 1$
just count
obs i in cluster
 k if $r_{ik} > 1$
 $= N_k$ ✓

$$\hat{\mu}_k = \frac{1}{N_k^{\text{soft}}} \sum_{i=1}^N r_{ik} x_i$$

only add x_i if i in k
($r_{ik} > 1$)

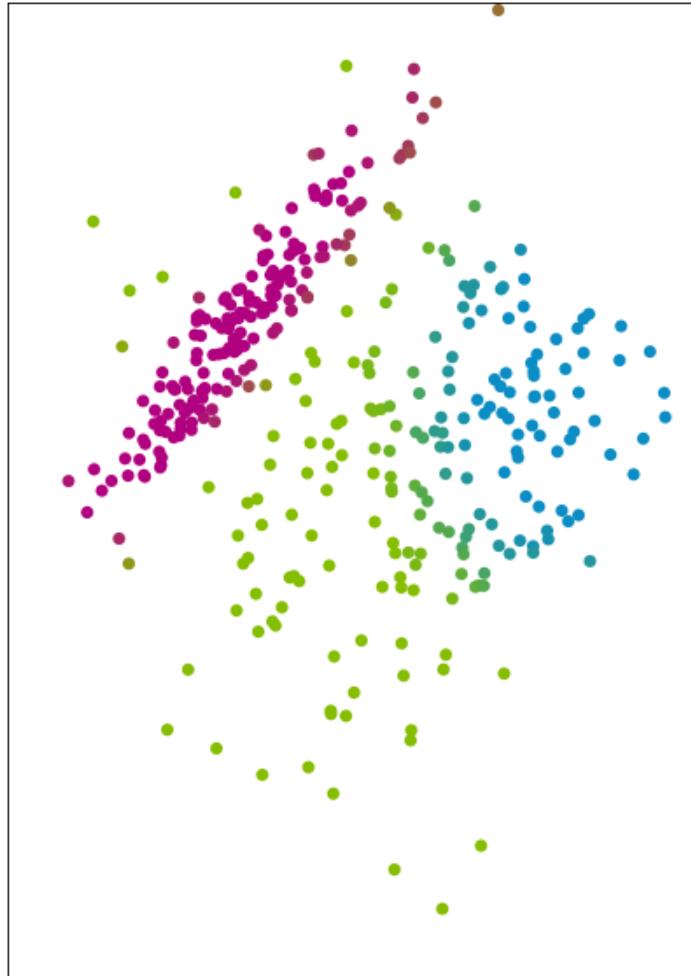
$$\rightarrow \frac{1}{N_k} \sum_{i \in k} x_i$$
 ✓

$$\hat{\Sigma}_k = \frac{1}{N_k^{\text{soft}}} \sum_{i=1}^N r_{ik} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T$$

same as above

$$= \frac{1}{N_k} \sum_{i \in k} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T$$

Part 2b summary



Still straightforward
to compute cluster
parameter estimates
from soft assignments

Expectation maximization (EM)

Expectation maximization (EM): An iterative algorithm

Motivates an iterative algorithm:


1. **E-step:** estimate cluster responsibilities
given current parameter estimates

$$\hat{r}_{ik} = \frac{\hat{\pi}_k N(x_i | \hat{\mu}_k, \hat{\Sigma}_k)}{\sum_{j=1}^K \hat{\pi}_j N(x_i | \hat{\mu}_j, \hat{\Sigma}_j)}$$

2. **M-step:** maximize likelihood over
parameters given current responsibilities

$$\hat{\pi}_k, \hat{\mu}_k, \hat{\Sigma}_k | \{\hat{r}_{ik}, x_i\}$$

EM for mixtures of Gaussians in pictures – initialization




Initialize
iter counter
 $\{\pi_k^{(0)}, \mu_k^{(0)}, \hat{\Sigma}_k^{(0)}\}$

Then compute

$$\hat{r}_i^{(1)}$$

$$\hat{r}_i^{(1)} = [0.52 \quad 0.4 \quad 0.08] \quad \text{fuchsia} \quad \text{blue} \quad \text{green}$$

EM for mixtures of Gaussians in pictures – after 1st iteration




Maximize likelihood
given soft assign. $r_{ik}^{(1)}$

$$\rightarrow \{\hat{\pi}_k^{(1)}, \hat{\mu}_k^{(1)}, \hat{\Sigma}_k^{(1)}\}$$

Then recompute responsibilities

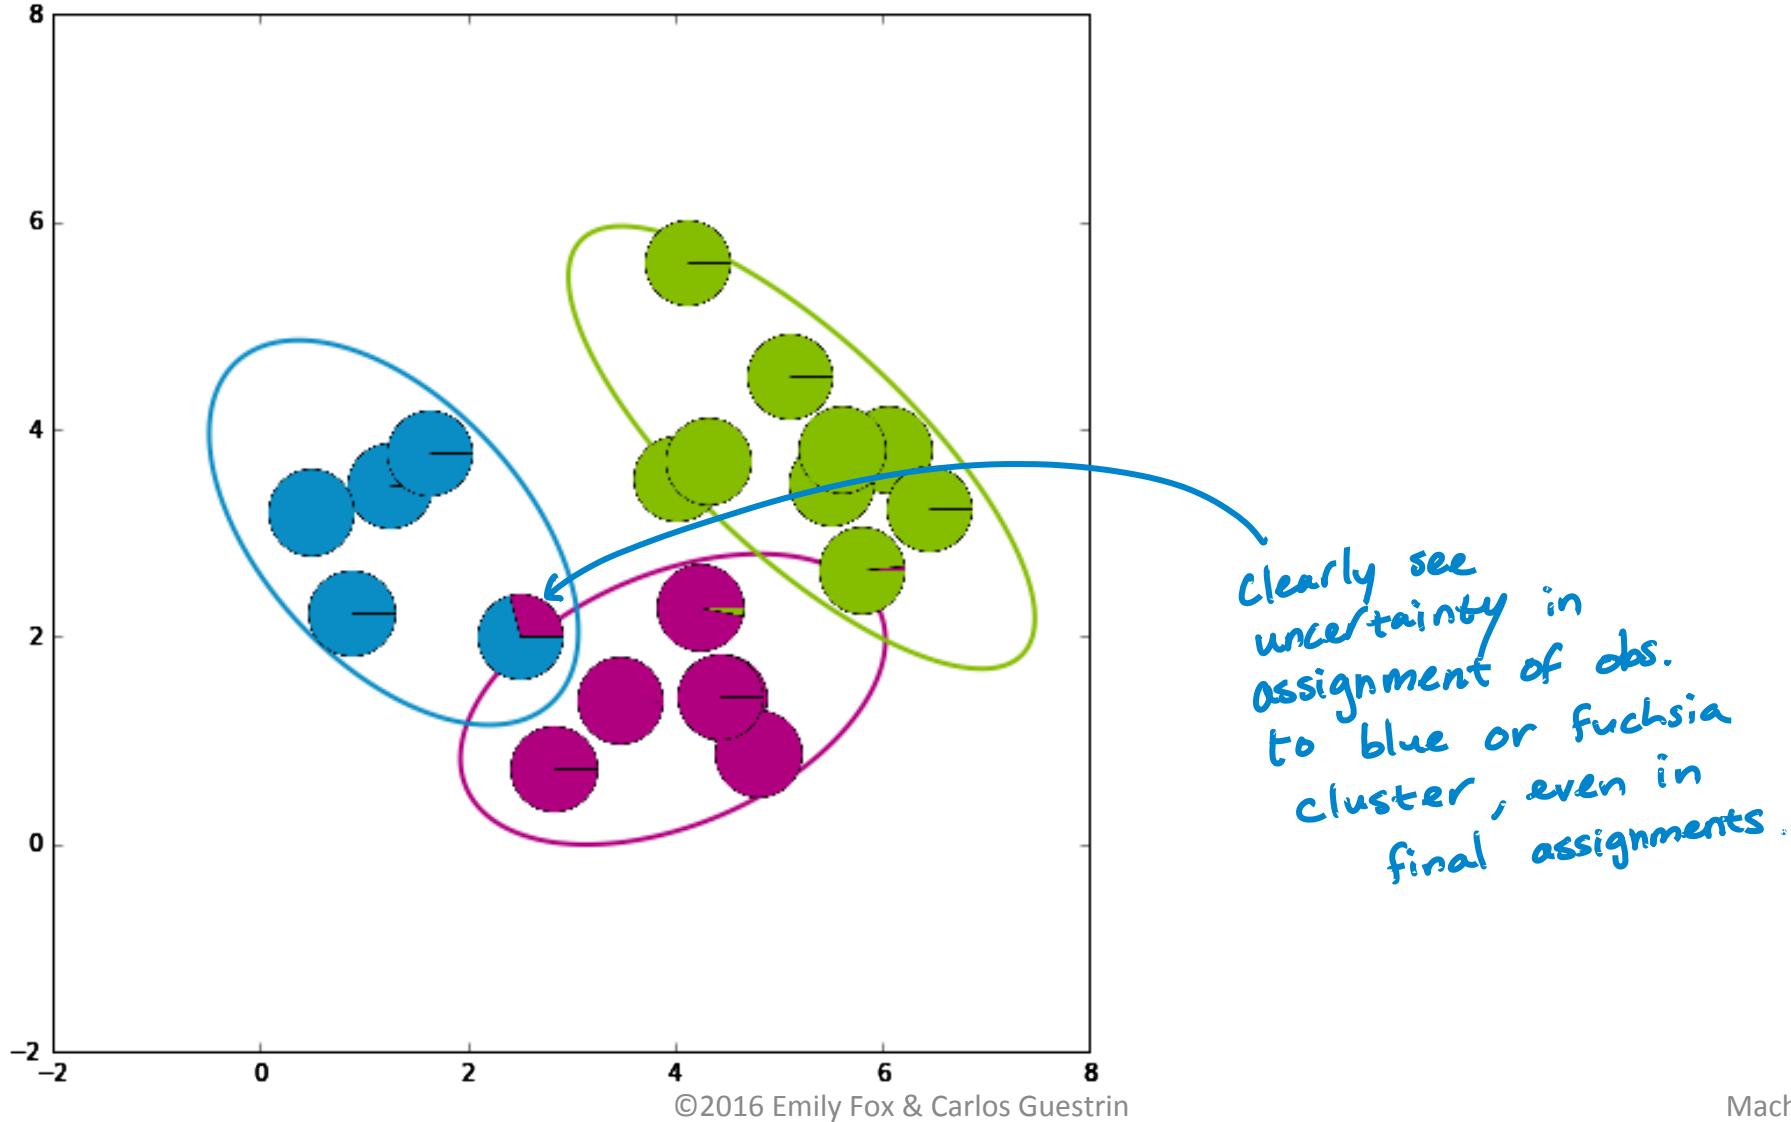
$$\hat{r}_{ik}^{(2)}$$

EM for mixtures of Gaussians in pictures – after 2nd iteration




rinse
+
repeat
until convergence

EM for mixtures of Gaussians in pictures – converged solution



EM for mixtures of Gaussians in pictures - replay



The nitty gritty of EM

Convergence of EM

- EM is a coordinate-ascent algorithm
 - Can equate E-and M-steps with alternating maximizations of an objective function
- Converges to a local mode
- We will assess via (log) likelihood of data under current parameter and responsibility estimates

Initialization


- Many ways to initialize the EM algorithm
- Important for convergence rates and quality of local mode found
- Examples:
 - Choose K observations at random to define K “centroids”. Assign other observations to nearest centroid to form initial parameter estimates.
 - Pick centers sequentially to provide good coverage of data like in k-means++
 - Initialize from k-means solution
 - Grow mixture model by splitting (and sometimes removing) clusters until K clusters are formed

Overfitting of MLE

Maximizing likelihood can **overfit to data**

Imagine at K=2 example with one obs assigned to cluster 1 and others assigned to **cluster 2**

- What parameter values maximize likelihood?



Set center equal to
point and shrink
variance to 0

Likelihood goes to ∞ !

Overfitting in high dims

Doc-clustering example:

Imagine only 1 doc assigned to cluster k has word w
(or all docs in cluster agree on count of word w)

Likelihood maximized by setting $\mu_k[w] = \mathbf{x}_i[w]$ and $\sigma_{w,k}^2 = 0$

Likelihood of any doc with different count on word w being in cluster k is 0!

Simple regularization of M-step for mixtures of Gaussians

Simple fix: Don't let variances $\rightarrow 0!$

Add small amount to diagonal of covariance estimate

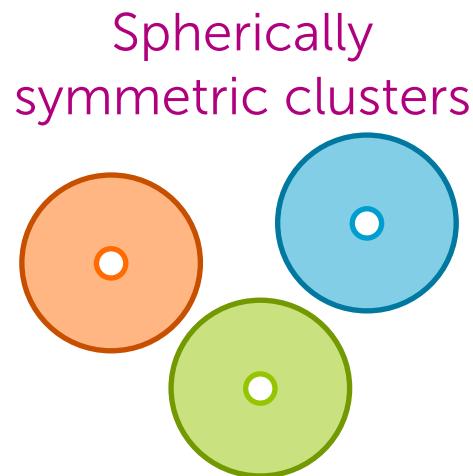
Alternatively, take Bayesian approach and place prior on parameters.

Similar idea, but all parameter estimates are “smoothed” via cluster pseudo-observations.

Relationship to k-means

Consider Gaussian mixture model with

$$\Sigma = \begin{pmatrix} \sigma^2 & & & \\ & \sigma^2 & & \\ & & \sigma^2 & \\ & \ddots & & \ddots & \\ & & & & \sigma^2 \end{pmatrix}$$



and let the variance parameter $\sigma \rightarrow 0$

Datapoint gets fully assigned to nearest center, just as in k-means

- Spherical clusters with equal variances, so **relative likelihoods** just function of distance to cluster center
- As variances $\rightarrow 0$, **likelihood ratio becomes 0 or 1**
- Responsibilities weigh in cluster proportions, but dominated by likelihood disparity

$$\hat{r}_{ik} = \frac{\hat{\pi}_k N(x_i | \hat{\mu}_k, \sigma^2 I)}{\sum_{j=1}^K \hat{\pi}_j N(x_i | \hat{\mu}_j, \sigma^2 I)}$$

Infinitesimally small variance EM = k-means

1. **E-step:** estimate cluster responsibilities given current parameter estimates

$$\hat{r}_{ik} = \frac{\hat{\pi}_k N(x_i | \hat{\mu}_k, \sigma^2 I)}{\sum_{j=1}^K \hat{\pi}_j N(x_i | \hat{\mu}_j, \sigma^2 I)} \in \{0, 1\}$$

Infinitesimally small 

Decision based on distance to nearest cluster center 

2. **M-step:** maximize likelihood over parameters given current responsibilities (**hard assignments!**)

$$\hat{\pi}_k, \hat{\mu}_k \mid \{\hat{r}_{ik}, x_i\}$$

Summary for mixture models and the EM algorithm

What you can do now...

- Interpret a probabilistic model-based approach to clustering using mixture models
- Describe model parameters
- Motivate the utility of soft assignments and describe what they represent
- Discuss issues related to how the number of parameters grow with the number of dimensions
 - Interpret diagonal covariance versions of mixtures of Gaussians
- Compare and contrast mixtures of Gaussians and k-means
- Implement an EM algorithm for inferring soft assignments and cluster parameters
 - Determine an initialization strategy
 - Implement a variant that helps avoid overfitting issues